Quick Decision Rules:

If Question Says...

"Find all customers who placed orders"

"Find employees in departments 10, 20, 30"

"Get department name with every employee"

"Find employee whose salary = max salary"

"Compare each row to a group result"

EXISTS

Multi-Row Subquery (IN)

Use

JOIN

Single-Row Subquery

Correlated Subquery

Summary: When to Use What?

Concept	Use When	Typical Syntax / Use	Returns	Best Used For
• Single Row Subquery	Subquery returns only 1 value	=, >, <, etc.	One value	Comparing against max/min/avg, filtering
Multi Row Subquery	Subquery returns one column, multiple rows	IN, NOT IN, ANY, ALL	One column, many rows	Filtering based on list of values
Multi Row, Multi Column Subquery	Subquery returns multiple columns and rows	IN (SELECT col1, col2), used with row constructors	Multiple columns & rows	Row-wise filtering, composite conditions
• Correlated Subquery	Subquery depends on outer row	Subquery inside WHERE or SELECT, uses outer table alias	Varies (evaluated per row)	Per-row logic like group-wise comparison
• INNER JOIN	You need matching rows from 2+ tables	SELECT FROM A JOIN B ON A.x = B.x	Combined rows	Fetching related data
• LEFT/RIGHT JOIN	You need non- matching rows too	LEFT JOIN, RIGHT JOIN	Includes NULLs for missing match	Optional relationships, reporting all even if unmatched
• EXISTS / NOT EXISTS	You want to check presence/absence	WHERE EXISTS (SELECT)	Boolean logic	Presence, efficient for large data

Tip for Your Students:

- Use **IN** for simple matching
- Use joins when you need columns from both tables
- Use correlated subqueries when logic changes per row
- Use **EXISTS** when just checking presence is enough (faster than **IN** on big data)

✓ 1. "Find employees who earn more than avg salary of their dept"

Use: Correlated Subquery

Because average salary changes per department, you need a subquery that depends on each employee's dept_id.

```
SELECT e.employee_id, e.name, e.salary, e.dept_id
FROM employees e
WHERE e.salary > (
    SELECT AVG(e2.salary)
    FROM employees e2
    WHERE e2.dept_id = e.dept_id
);
```

✓ This compares each employee's salary with the **average salary of their department**.

List employees **who earn more than the average salary** of their own department.

Sample employees table:

employee_id name salary dept_id

101	кај	60,000	10
102	Priya	50,000	10
103	Aarav	40,000	10
104	Neha	70,000	20
105	Rohan	60,000	20
106	Kavya	90,000	30

Step-by-step Dry Run:

• For Raj (Dept 10, Salary 60,000):

Subquery:

```
SELECT AVG(salary) FROM employees e2 WHERE dept_id = 10 \rightarrow (60,000 + 50,000 + 40,000) / 3 = 50,000 Check: 60,000 > 50,000 \checkmark \rightarrow Include Raj
```

• For Priya (Dept 10, Salary 50,000):

Same AVG: 50,000

Check: 50,000 > 50,000 \times \rightarrow Exclude

• For Aarav (Dept 10, Salary 40,000):

Check: $40,000 > 50,000 \times \to \text{Exclude}$

• For Neha (Dept 20, Salary 70,000):

Subquery:

SELECT AVG(salary) FROM employees WHERE dept_id = 20

 \rightarrow (70,000 + 60,000) / 2 = 65,000

Check: 70,000 > 65,000 ✓ → Include Neha

• For Rohan (Dept 20, Salary 60,000):

Check: $60,000 > 65,000 \times \to Exclude$

• For Kavya (Dept 30, Salary 90,000):

Only one person in Dept $30 \rightarrow AVG = 90,000$ Check: $90,000 > 90,000 \times \rightarrow Exclude$

Final Output:

employee_id name salary dept_id

101 Raj 60,000 10 104 Neha 70,000 20

2. "Find all customers who placed orders"

Use: EXISTS

• You just need to check if a matching order **exists**, not fetch it.

```
SELECT c.customer_id, c.name
FROM customers c
WHERE EXISTS (
    SELECT 1
    FROM orders o
    WHERE o.customer_id = c.customer_id
);
```

EXISTS is faster than IN on large datasets and ignores duplicates.

Dry Run (How Oracle Executes It):

Suppose:

customers table:

```
customer_idname101Alice102Bob103Carol
```

orders table:

```
        order_id
        customer_id

        1
        101

        2
        103
```

Step-by-step execution:

```
1. Pick customer Alice (101)
```

```
\rightarrow Run subquery: SELECT 1 FROM orders o WHERE o.customer_id = 101 \rightarrow Result: Exists \checkmark \rightarrow Include Alice
```

2. Pick customer Bob (102)

```
→ Run subquery: SELECT 1 FROM orders o WHERE o.customer_id = 102
→ Result: No match X → Exclude Bob
```

3. Pick customer Carol (103)

```
\rightarrow Run subquery: SELECT 1 FROM orders o WHERE o.customer_id = 103
```

```
ightarrow Result: Exists 
ightharpoonup 
ightharpoonup Include Carol
```

☑ Final Output:

```
customer_id name
101 Alice
103 Carol
```

© Query (Correlated Subquery)

```
SELECT e.employee_id, e.name, e.salary, e.dept_id
FROM employees e
WHERE e.salary > (
    SELECT AVG(e2.salary)
    FROM employees e2
    WHERE e2.dept_id = e.dept_id
);
```

Query Goal:

Select employees who earn **more than the average salary of their own department**.

Sample employees Table

employee_idnamesalarydept_id101Raj60,00010102Priya50,00010103Aarav40,00010104Neha70,00020105Rohan60,00020106Kavya90,00030

Dry Run – Step-by-Step Execution:

We will process **each row** and compare its salary with the **average salary of its department**.

Raj (Dept 10, Salary 60,000)

```
• Subquery:
```

```
SELECT AVG(salary) FROM employees WHERE dept_id = 10 \rightarrow (60,000 + 50,000 + 40,000) / 3 = 50,000
```

- Condition: 60,000 > 50,000 <
 - → Include Raj

Priya (Dept 10, Salary 50,000)

- Avg = 50,000
- Condition: 50,000 > 50,000 X
 - → Exclude Priya

🔽 Aarav (Dept 10, Salary 40,000)

- Avg = 50,000
- $40,000 > 50,000 \times$
 - → Exclude Aarav

🔽 Neha (Dept 20, Salary 70,000)

- Avg of Dept 20 = (70,000 + 60,000) / 2 = 65,000
- 70,000 > 65,000
 - → Include Neha

🔽 Rohan (Dept 20, Salary 60,000)

- Avg = 65,000
- $60,000 > 65,000 \times$
 - → Exclude Rohan

Kavya (Dept 30, Salary 90,000)

- Only employee in dept 30
- Avg = 90,000
- 90,000 > 90,000 X
 - → Exclude Kavya

Final Output:

employee_id name salary dept_id

101 Raj 60000 10 104 Neha 70000 20

Trainer Notes:

- This query is a **classic use of correlated subquery**: every row runs its **own AVG()** query.
- It's very useful to **highlight per-group comparisons** without using JOIN or GROUP BY.

Would you like the same dry-run shown in .pptx or .sql format with visual callouts for class?



List employees who earn more than the average salary of their department, using JOIN + **GROUP BY** instead of a correlated subquery.

Method 1: Using a Derived Table (Subquery + JOIN)

```
sql
CopyEdit
SELECT e.employee_id, e.name, e.salary, e.dept_id
FROM employees e
JOIN (
   SELECT dept_id, AVG(salary) AS avg_salary
   FROM employees
   GROUP BY dept_id
) dept_avg ON e.dept_id = dept_avg.dept_id
WHERE e.salary > dept_avg.avg_salary;
```

Explanation:

1. **Inner subquery**: Calculates average salary per department.

```
SELECT dept_id, AVG(salary)
FROM employees
GROUP BY dept_id
```

- 2. **JOIN**: Matches each employee with their department's average.
- 3. **WHERE clause**: Filters employees whose salary > average.

🔑 Result

employee_id name salary dept_id

60000 10 101 Raj 104 Neha 70000 20

✓ **Same result** as with the correlated subquery — but may **perform better** on large datasets.

Bring a list of all those employees who earn more than the average salary of their department, but don't use a correlated subquery.



```
CopyEdit
SELECT e.employee_id, e.name, e.salary, e.dept_id
FROM employees e
JOIN (
    SELECT dept_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY dept_id
) dept_avg
ON e.dept_id = dept_avg.dept_id
WHERE e.salary > dept_avg.avg_salary;
```

Q Breakdown (Simple Explanation):

Step 1: Inner Query – Group By Department

```
sql
CopyEdit
SELECT dept_id, AVG(salary) AS avg_salary
FROM employees
GROUP BY dept_id;
```

This query calculates the average salary of each department.

Example output:

dept_id avg_salary

10 50,000 20 65,000 30 90,000

Step 2: JOIN Outer Table (employees e)

```
sql
CopyEdit
FROM employees e
JOIN (...) dept_avg
ON e.dept_id = dept_avg.dept_id
```

Now we are **joining each employee** with the **average salary of their department**, using dept_id as the matching point.

Step 3: Filter Employees Whose Salary > Avg Salary

```
sql
CopyEdit
WHERE e.salary > dept_avg.avg_salary;
```

Here we are only showing those employees whose salary is more than the average salary of their own department.

Sample Data:

employee_id name salary dept_id

```
101 Raj 60,000 10

102 Priya 50,000 10

103 Aarav 40,000 10

104 Neha 70,000 20

105 Rohan 60,000 20

106 Kavya 90,000 30
```

Avg salaries from subquery:

- Dept $10 \rightarrow (60+50+40)/3 = 50,000$
- Dept 20 \rightarrow (70+60)/2 = 65,000
- Dept $30 \rightarrow 90,000$

Final Output:

employee_id name salary dept_id

101 Raj 60,000 10 104 Neha 70,000 20

dept_avg is an alias — not a real table.

Q Explanation:

In this part of the query:

```
JOIN (
   SELECT dept_id, AVG(salary) AS avg_salary
   FROM employees
   GROUP BY dept_id
) dept_avg
```

- You are writing a **subquery** (also called a **derived table** or **inline view**).
- dept_avg is simply a **nickname** (alias) for that subquery's result.
- Think of it like: JOIN (result_set) AS dept_avg

Analogy

it's like you made a temporary result with:

```
dept_id | avg_salary
-----10 | 50000
```

20 | 65000

And gave that result a $nickname \rightarrow dept_avg$, which you can now use like a table:

dept_avg.dept_id
dept_avg.avg_salary

☑ 3. "Find employees in departments 10, 20, 30"

Use: Multi-Row Subquery (IN)

```
You're filtering based on a list of values (can be static or dynamic).
```

Simple value list filter — one column, multiple rows.

4. "Get department name with every employee"

Use: INNER JOIN

Nou need columns from both tables: employee and department.

```
sql
CopyEdit
SELECT e.employee_id, e.name, d.department_name
FROM employees e
JOIN departments d ON e.dept_id = d.dept_id;
```

✓ JOIN is the best when you're **combining columns** from multiple tables.

5. "Find employee whose salary = max salary"

Use: Single-Row Subquery

• Only one max salary value, so subquery returns **one row**.

```
sql
CopyEdit
SELECT *
FROM employees
WHERE salary = (
     SELECT MAX(salary)
     FROM employees
);
```

✓ Use = with **scalar subquery** — only works if inner query returns exactly **one value**.

6. "Compare each row to a group result"

Use: Correlated Subquery (again)

Same logic as #1, e.g., compare to **max salary of their dept**:

```
sql
CopyEdit
SELECT e.employee_id, e.name, e.salary
FROM employees e
WHERE salary = (
    SELECT MAX(salary)
    FROM employees e2
    WHERE e2.dept_id = e.dept_id
);
```

Every row gets its own inner query result.



Bonus Tip:

If you just want to check condition, use:

- ✓ IN for simple list matching
- **EXISTS** for **faster lookup** and avoiding NULL problems
- Correlated subqueries for row-wise logic
- **✓** Joins for **data from multiple tables**